

Sustainable Smart City Assistant using IBM Granite LLM

1. Introduction

Project Title: Sustainable Smart City Assistant using IBM Granite LLM

Team Member 1: Madhumathi.S

Team Member 2: Kokila.K

2. Project Overview

Purpose: The purpose of this project is to design an AI-powered assistant that helps cities transition into sustainable, intelligent, and citizen-friendly ecosystems. Powered by IBM Granite LLM within the Watsonx platform, the assistant enables cities to manage resources efficiently, engage with communities, and support officials in strategic decision-making.

It integrates AI-driven forecasting, anomaly detection, and policy summarization to reduce environmental impact and enhance governance transparency.

Features:

- Conversational Interface: Natural language interaction for citizens and officials.
- Policy Summarization: Simplifies lengthy government policies into clear, actionable insights.
- Resource Forecasting: Predicts future trends in water, energy, and waste usage.
- Eco-Tip Generator: Provides personalized sustainability recommendations.
- Citizen Feedback Loop: Collects, analyzes, and integrates public input.
- KPI Forecasting: Assists city officials with strategic planning and performance monitoring.
- Anomaly Detection: Identifies unusual patterns in urban data for early issue resolution.
- Multimodal Input: Supports text, PDFs, and CSVs for analysis.
- Streamlit/Gradio UI: Intuitive interface for both citizens and officials.

3. Architecture

Frontend (Streamlit/Gradio): Provides an interactive dashboard with modules for chat, file uploads, reports, and forecasting.

Backend (FastAPI): REST API framework handling requests for summarization, forecasting, eco-tips, and anomaly detection.

LLM Integration (IBM Watsonx Granite): Granite models enable natural language understanding, summarization, and advice generation.

Vector Search (Pinecone): Stores embedded documents and enables semantic search of policies and citizen queries.

ML Modules (Scikit-learn, Pandas, Matplotlib): Implements forecasting and anomaly detection using time-series data.

4. Setup Instructions

Prerequisites:

- Python 3.9+
- pip and virtual environment tools
- IBM Watsonx API key

- Pinecone API key
- Internet connection

Installation Process:

1. Clone repository
2. Install dependencies from requirements.txt
3. Configure .env with credentials
4. Start FastAPI backend
5. Run Streamlit dashboard
6. Upload data and interact with modules

5. Folder Structure

app/ – FastAPI backend (routes, models, integration)

app/api/ – Modular API routes (chat, feedback, reports, embeddings)

ui/ – Frontend Streamlit components

smart_dashboard.py – Main entry for Streamlit dashboard

granite_llm.py – Handles IBM Granite LLM communication

document_embedder.py – Document embedding & Pinecone storage

kpi_file_forecaster.py – Forecasting trends in urban data

anomaly_file_checker.py – Detects anomalies in datasets

report_generator.py – AI-generated sustainability reports

6. Running the Application

- Launch FastAPI server
- Start Streamlit UI
- Navigate via sidebar to access dashboards, chat, eco-tips, forecasting, and reports
- Upload files (PDF/CSV) and interact with the AI assistant
- Get real-time summaries, insights, and predictions

7. API Documentation

POST /chat/ask – User queries, AI-generated responses

POST /upload-doc – Uploads and embeds documents in Pinecone

GET /search-docs – Returns similar policies using semantic search

GET /get-eco-tips – Provides sustainability suggestions

POST /submit-feedback – Stores citizen feedback

8. Authentication

- Token-based authentication (JWT/API Keys)
- IBM Cloud OAuth2 credentials
- Role-based access (citizen, official, admin)

- Planned: User sessions & history tracking

9. User Interface

- Sidebar navigation
- KPI dashboards with visual summaries
- Tabbed layouts (Chat, Eco Tips, Forecasting, Reports)
- Real-time feedback forms
- Exportable PDF reports

10. Testing

- Unit Testing – For prompt functions & utilities
- API Testing – Swagger UI/Postman
- Manual Testing – File uploads, chat outputs, report accuracy
- Edge Case Handling – Large files, malformed input, invalid keys

11. Screenshots

(To be added after implementation – dashboards, chat, forecasting graphs, eco-tips, reports)

12. Known Issues

- Requires stable internet for IBM Watsonx API calls
- Large file uploads may increase processing time
- Limited offline capability

13. Future Enhancements

- Multilingual support
- Mobile app version
- Integration with IoT smart city sensors
- Real-time traffic & pollution monitoring
- AI-driven citizen engagement dashboards